

P8106_hw4_yh3554

Data Science II Homework 4

Yi Huang

2023-04-21

Contents

Problem 1.	2
Data cleaning	3
(a)	4
(i) Build a regression tree on train data	4
(ii) create a plot of the tree	5
(b)	6
(i) Perform Random forest on train data	6
(ii) Report the variable importance and the test error.	7
(c)	7
(i) Perform Boosting on the training data	7
(ii) Report the variable importance and the test error.	8
Problem 2.	9
Data cleaning	9
(a)	10
(i)	10
(ii)	12
(b)	15
(i)	15
(ii)	17
Additional analysis: comparing classification tree and bootstrap	17

```
library(tidyverse)
library(dplyr)
library(knitr)
library(caret)
library(ISLR)
library(mlbench)
library(rpart)
library(rpart.plot)
library(party)
library(partykit)
library(pROC)
library(ranger)
library(gbm)
library(pdp)
library(ggplot2)
library(parallel)
library(doParallel)
```

Problem 1.

In this exercise, we will build tree-based models using the College data (see “College.csv” in Homework 2). The response variable is the out-of-state tuition (Outstate). Partition the dataset into two parts: training data (80%) and test data (20%).

The predictors are:

- Apps: Number of applications received
- Accept: Number of applications accepted
- Enroll: Number of new students enrolled
- Top10perc: Pct. new students from top 10% of H.S. class
- Top25perc: Pct. new students from top 25% of H.S. class
- F.Undergrad: Number of fulltime undergraduates
- P.Undergrad: Number of parttime undergraduates
- Room.Board: Room and board costs
- Books: Estimated book costs
- Personal: Estimated personal spending
- PhD: Pct. of faculty with Ph.D.’s
- Terminal: Pct. of faculty with terminal degree
- S.F.Ratio: Student/faculty ratio

- perc.alumni: Pct. alumni who donate
- Expend: Instructional expenditure per student
- Grad.Rate: Graduation rate

Data cleaning

```
# load data
dat <- read.csv("data/College.csv")[, -1]
dat <- na.omit(dat)
head(dat)
```

```
##   Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad Outstate
## 1 1660  1232   721      23        52      2885      537      7440
## 2 2186  1924   512      16        29      2683     1227     12280
## 3 1428  1097   336      22        50      1036      99      11250
## 4  417   349   137      60        89       510      63      12960
## 5  193   146    55      16        44       249     869      7560
## 6  587   479   158      38        62       678      41     13500
##   Room.Board Books Personal PhD Terminal S.F.Ratio perc.alumni Expend Grad.Rate
## 1      3300   450    2200  70      78      18.1      12    7041      60
## 2      6450   750    1500  29      30      12.2      16   10527      56
## 3      3750   400    1165  53      66      12.9      30    8735      54
## 4      5450   450     875  92      97       7.7      37   19016      59
## 5      4120   800    1500  76      72      11.9       2   10922      15
## 6      3335   500     675  67      73       9.4      11    9727      55
```

```
summary(dat)
```

```
##           Apps           Accept           Enroll           Top10perc
##  Min.   :  81   Min.   :  72   Min.   : 35.0   Min.   : 1.00
## 1st Qu.: 619   1st Qu.: 501   1st Qu.: 206.0   1st Qu.: 17.00
## Median : 1133   Median : 859   Median : 328.0   Median : 25.00
## Mean   : 1978   Mean   : 1306   Mean   : 456.9   Mean   : 29.33
## 3rd Qu.: 2186   3rd Qu.: 1580   3rd Qu.: 520.0   3rd Qu.: 36.00
## Max.   :20192   Max.   :13007   Max.   :4615.0   Max.   :96.00
##   Top25perc   F.Undergrad   P.Undergrad   Outstate
##  Min.   : 9.00   Min.   : 139   Min.   : 1   Min.   : 2340
## 1st Qu.: 42.00   1st Qu.: 840   1st Qu.: 63   1st Qu.: 9100
## Median : 55.00   Median : 1274   Median : 207   Median : 11200
## Mean   : 56.96   Mean   : 1872   Mean   : 434   Mean   : 11802
## 3rd Qu.: 70.00   3rd Qu.: 2018   3rd Qu.: 541   3rd Qu.: 13970
## Max.   :100.00   Max.   :27378   Max.   :10221   Max.   :21700
##   Room.Board   Books   Personal   PhD
##  Min.   :2370   Min.   : 250.0   Min.   : 250   Min.   : 8.00
## 1st Qu.:3736   1st Qu.: 450.0   1st Qu.: 800   1st Qu.: 60.00
## Median :4400   Median : 500.0   Median : 1100   Median : 73.00
## Mean   :4586   Mean   : 547.5   Mean   : 1214   Mean   : 71.09
## 3rd Qu.:5400   3rd Qu.: 600.0   3rd Qu.: 1500   3rd Qu.: 85.00
## Max.   :8124   Max.   :2340.0   Max.   :6800   Max.   :100.00
##   Terminal   S.F.Ratio   perc.alumni   Expend   Grad.Rate
##  Min.   : 24.00   Min.   : 2.50   Min.   : 2.00   Min.   : 3186   Min.   : 15
```

```
## 1st Qu.: 68.00    1st Qu.:11.10    1st Qu.:16.00    1st Qu.: 7477    1st Qu.: 58
## Median : 81.00    Median :12.70    Median :25.00    Median : 8954    Median : 69
## Mean   : 78.53    Mean   :12.95    Mean   :25.89    Mean   :10486    Mean   : 69
## 3rd Qu.: 92.00    3rd Qu.:14.50    3rd Qu.:34.00    3rd Qu.:11625    3rd Qu.: 81
## Max.   :100.00    Max.    :39.80    Max.    :64.00    Max.    :56233    Max.    :118
```

```
set.seed(123)
train_rows <- createDataPartition(y = dat$Outstate,
                                   p = 0.8,
                                   list = FALSE)

# training data
dat_train <- dat[train_rows, ]
x <- dat_train %>% select(-Outstate)
y <- dat_train$Outstate
# test data
dat_test <- dat[-train_rows, ]
x2 <- dat_test %>% select(-Outstate)
y2 <- dat_test$Outstate

set.seed(123)
# resampling method
ctrl <- trainControl(method = "cv")
```


(a)

Build a regression tree on the training data to predict the response. Create a plot of the tree.

(i) Build a regression tree on train data

```
set.seed(123)
rpart.fit <- train(Outstate ~ . ,
                  dat_train,
                  method = "rpart",
                  tuneGrid = data.frame(cp = exp(seq(-6,1, length = 100))),
                  trControl = ctrl)

ggplot(rpart.fit, highlight = TRUE)
```

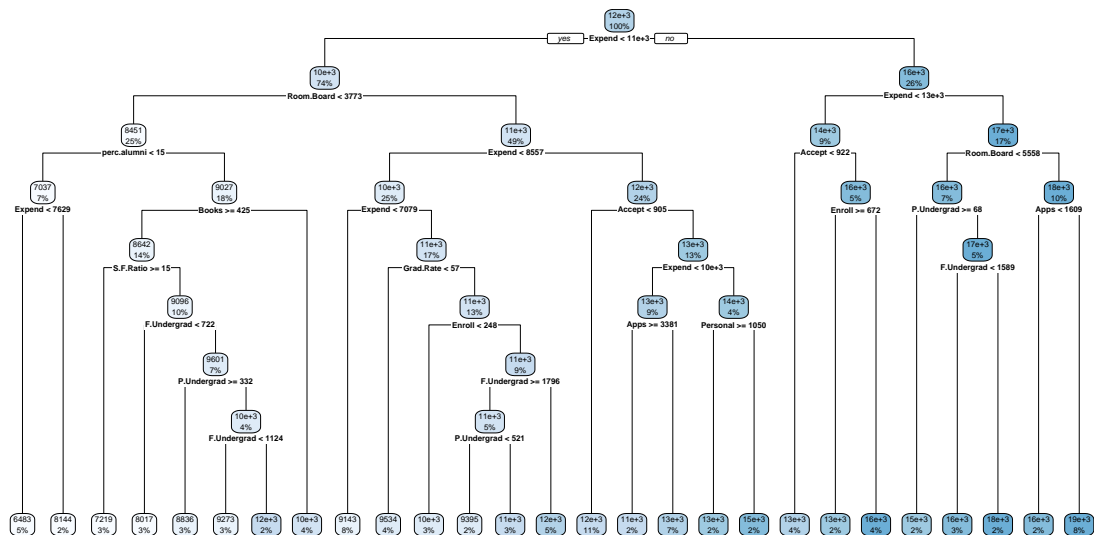


```
rpart.fit$finalModel$tuneValue[[1]]
```

```
## [1] 0.002478752
```

(ii) create a plot of the tree

```
rpart.plot(rpart.fit$finalModel)
```




(b)

Perform random forest on the training data. Report the variable importance and the test error.

(i) Perform Random forest on train data

```
rf.grid <- expand.grid(mtry = 1:16,  
                      splitrule = "variance",  
                      min.node.size = 1:6)  
  
set.seed(123)  
no_cores <- detectCores() - 1  
cl <- makePSOCKcluster(no_cores)  
registerDoParallel(cl)  
rf.fit <- train(Outstate ~ . ,  
               dat_train,  
               method = "ranger",  
               tuneGrid = rf.grid,  
               trControl = ctrl)  
  
stopCluster(cl)  
registerDoSEQ()  
ggplot(rf.fit, highlight = TRUE)
```




```
rf.fit$bestTune
```

```
## mtry splitrule min.node.size  
## 44 8 variance 2
```

The best tuning parameters are `mtry = 8` with `min.node.size = 2`.

(ii) Report the variable importance and the test error.

```
set.seed(123)
# variable importance
rf.final.per <- ranger(Outstate ~ . ,
                      dat_train,
                      mtry = rf.fit$bestTune[[1]],
                      splitrule = "variance",
                      min.node.size = rf.fit$bestTune[[3]],
                      importance = "permutation",
                      scale.permutation.importance = TRUE)
barplot(sort(ranger::importance(rf.final.per), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



```
# test error
rf.predict <- predict(rf.fit, newdata = dat_test)
rf.RMSE <- RMSE(rf.predict, y2)
rf.RMSE
```

```
## [1] 1850.927
```

The top 6 most important variables are Expend, Room.Board, Apps, Accept, Terminal, and Top10perc. The RMSE of test set is 1850.93.

(c)

Perform boosting on the training data. Report the variable importance and the test error.


(i) Perform Boosting on the training data

```
gbm.grid <- expand.grid(n.trees = c(1000, 2000, 3000, 4000, 5000),
                      interaction.depth = 1:6,
                      shrinkage = seq(0.001, 0.005, by = 0.002),
                      n.minobsinnode = c(3:15))

set.seed(123)
```

```
no_cores <- detectCores() - 1
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
gbm.fit <- train(Outstate ~ . ,
  dat_train,
  method = "gbm",
  tuneGrid = gbm.grid,
  trControl = ctrl,
  verbose = FALSE)

stopCluster(cl)
registerDoSEQ()
ggplot(gbm.fit, highlight = TRUE)
```



```
gbm.fit$bestTune
```

```
##      n.trees interaction.depth shrinkage n.minobsinnode
## 1166      1000              6      0.005             15
```

The best tuning parameters are `n.trees = 1000`, `interaction.depth = 6`, `shrinkage = 0.005` and `nminobsinode = 15`.

(ii) Report the variable importance and the test error.


```
set.seed(123)
gbm.final.per <- ranger(Outstate ~ . ,
  dat_train,
  n.trees = gbm.fit$bestTune[[1]],
```



```

splitrule = "variance",
interaction.depth = gbm.fit$bestTune[[2]],
shrinkage = gbm.fit$bestTune[[3]],
n.minobsinnode = gbm.fit$bestTune[[4]],
importance = "permutation",
scale.permutation.importance = TRUE)
barplot(sort(ranger::importance(gbm.final.per), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19))

```



```

# test error
gbm.predict <- predict(gbm.fit, newdata = dat_test)
gbm.RMSE <- RMSE(gbm.predict, y2)
gbm.RMSE

```

```
## [1] 1733.132
```

The top 6 most important variables are Expend, Room.Board, Apps, Terminal, Top10perc, and Accept. The RMSE of test set is 1733.13.

Problem 2.

This problem involves the OJ data in the ISLR package. The data contains 1070 purchases where the customers either purchased Citrus Hill or Minute Maid Orange Juice. A number of characteristics of customers and products are recorded. Create a training set containing a random sample of 700 observations, and a test set containing the remaining observations.

Data cleaning

```

data(OJ)
OJ <- na.omit(OJ)
set.seed(123)

train_rows2 <- createDataPartition(y = OJ$Purchase,
                                    p = 0.8,

```

```

                                list = FALSE)

# training data
OJ_train <- OJ[train_rows2, ]
# test data
OJ_test <- OJ[-train_rows2, ]

# resampling method
ctrl1 <- trainControl(method = "cv",
                      classProbs = TRUE)

set.seed(123)
ctrl2 <- trainControl(method = "cv",
                      classProbs = TRUE,
                      summaryFunction = twoClassSummary)

```

(a)


(i)

Build a classification tree using the training data, with Purchase as the response and the other variables as predictors. Which tree size corresponds to the lowest cross-validation error?

```

set.seed(123)
rpart.fit.OJ <- train(Purchase ~ . ,
                     OJ_train,
                     method = "rpart",
                     tuneGrid = data.frame(cp = exp(seq(-6,-3, len = 50))),
                     trControl = ctrl2,
                     metric = "ROC")
ggplot(rpart.fit.OJ, highlight = TRUE)

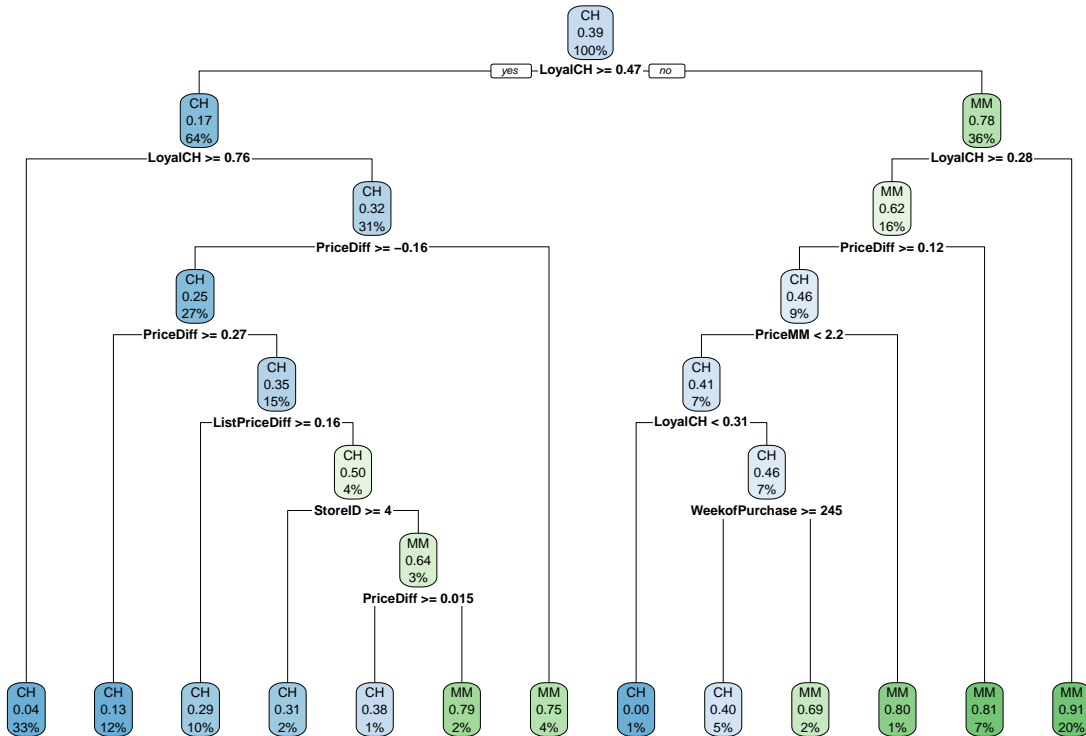
```



```
rpart.fit.OJ$bestTune
```

```
##          cp
## 11 0.004572226
```

```
# summary(rpart.fit.OJ)
rpart.plot(rpart.fit.OJ$finalModel)
```



The tree size of 13 has lowest cross-validation error with $cp = 0.004572$.

Note: tree size = number of split + 1

(ii)


Is this the same as the tree size obtained using the 1 SE rule?

```
set.seed(123)
tree1 <- rpart(formula = Purchase ~ . ,
               OJ_train,
               control = rpart.control(cp = 0))
cpTable <- printcp(tree1)

##
## Classification tree:
## rpart(formula = Purchase ~ ., data = OJ_train, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] ListPriceDiff LoyalCH PriceCH PriceDiff PriceMM
## [6] SpecialCH StoreID WeekofPurchase
##
## Root node error: 334/857 = 0.38973
##
## n= 857
##
## CP nsplit rel error xerror xstd
## 1 0.517964 0 1.00000 1.00000 0.042745
## 2 0.026946 1 0.48204 0.52096 0.035257
## 3 0.011976 3 0.42814 0.49401 0.034559
## 4 0.007485 6 0.39222 0.46108 0.033651
## 5 0.005988 8 0.37725 0.45808 0.033566
```

```
## 6 0.004491    12  0.35329 0.48503 0.034317
## 7 0.002994    14  0.34431 0.49701 0.034638
## 8 0.000998    16  0.33832 0.48802 0.034398
## 9 0.000000    19  0.33533 0.52395 0.035332
```

```
plotcp(tree1)
```



```
# rpart.plot(tree1)
```

```
set.seed(123)
```

```
# 1SE rule
```

```
minErr <- which.min(cpTable[,4])
```

```
tree2 <- prune(tree1,cp = cpTable[cpTable[,4]<cpTable[minErr,4]+cpTable[minErr,5],1][1])
```

```
cpTable <- printcp(tree2)
```

```
##
```

```
## Classification tree:
```

```
## rpart(formula = Purchase ~ ., data = OJ_train, control = rpart.control(cp = 0))
```

```
##
```

```
## Variables actually used in tree construction:
```

```
## [1] LoyalCH PriceDiff PriceMM
```

```
##
```

```
## Root node error: 334/857 = 0.38973
```

```
##
```

```
## n= 857
```

```
##
```

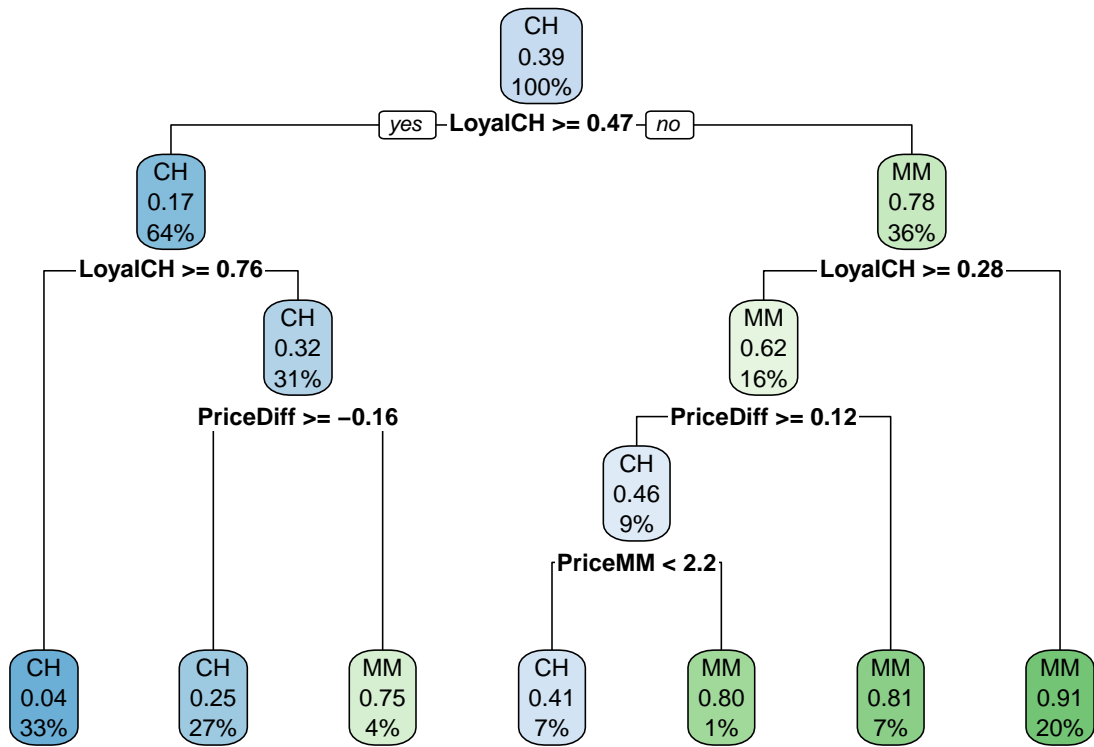
```
##          CP nsplit rel error  xerror    xstd
```

```
## 1 0.517964      0  1.00000 1.00000 0.042745
```


```
## 2 0.026946      1  0.48204 0.52096 0.035257
```

```
## 3 0.011976      3  0.42814 0.49401 0.034559
## 4 0.007485      6  0.39222 0.46108 0.033651
```

```
rpart.plot(tree2)
```



```
plotcp(tree2)
```



Under 1 SE rule, the tree size with lowest cross-validation error is 7. The tree size obtained by using cross validation is different from the tree size obtained by using 1 SE rule.

(b)


(i)

Perform boosting on the training data and report the variable importance.

```
gbmA.grid <- expand.grid(n.trees = c(1000,2000,3000,4000,5000),
                        interaction.depth = 1:6,
                        shrinkage = c(0.0005,0.001,0.002),
                        n.minobsinnode = 1)

set.seed(123)
no_cores <- detectCores() - 1
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
gbmA.fit <- train(Purchase ~ . ,
                  OJ_train,
                  tuneGrid = gbmA.grid,
                  trControl = ctrl2,
                  method = "gbm",
                  distribution = "adaboost",
                  metric = "ROC",
                  verbose = FALSE)

stopCluster(cl)
registerDoSEQ()
ggplot(gbmA.fit, highlight = TRUE)
```



```
summary(gbmA.fit$finalModel, las = 2, cBars = 19, cex.names = 0.6) %>%
  knitr::kable(digits = 3, caption = "Variable importance from boosting model")
```




Table 1: Variable importance from boosting model

	var	rel.inf
LoyalCH	LoyalCH	77.013
PriceDiff	PriceDiff	10.524
ListPriceDiff	ListPriceDiff	3.089
SalePriceMM	SalePriceMM	2.229
WeekofPurchase	WeekofPurchase	1.794
StoreID	StoreID	1.242
STORE	STORE	0.789
DiscCH	DiscCH	0.769
SpecialCH	SpecialCH	0.634
DiscMM	DiscMM	0.525
SalePriceCH	SalePriceCH	0.452
PriceMM	PriceMM	0.378
PriceCH	PriceCH	0.253
PctDiscCH	PctDiscCH	0.101
SpecialMM	SpecialMM	0.086
Store7Yes	Store7Yes	0.061
PctDiscMM	PctDiscMM	0.059

In the boosting model, the top 2 most important variables are `LoyalCH` and `PriceDiff`.

(ii)

What is the test error rate?

```
gbmA.pred <- predict(gbmA.fit, newdata = OJ_test, type = "raw")
error.rate.gbmA <- mean(gbmA.pred != OJ$Purchase[-train_rows2])
error.rate.gbmA
```

```
## [1] 0.1971831
```

The test error rate is 0.197.

Additional analysis: comparing classification tree and bootstrap

Report cross-validation results on train data

```
set.seed(123)
resamp <- resamples(list( ctrees = rpart.fit.OJ,
                        gbmA = gbmA.fit))

summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: ctrees, gbmA
## Number of resamples: 10
##
## ROC
##           Min.   1st Qu.   Median   Mean   3rd Qu.   Max. NA's
## ctrees 0.7777149 0.8547942 0.8879662 0.8794884 0.9177400 0.9268648    0
```

```
## gbmA    0.8546380 0.8899477 0.9073427 0.9055784 0.9278846 0.9519726    0
##
## Sens
##           Min.    1st Qu.    Median      Mean    3rd Qu.      Max. NA's
## ctree  0.7884615 0.8537736 0.8762700 0.8642598 0.8846154 0.9038462    0
## gbmA    0.8269231 0.8750000 0.9143687 0.8966255 0.9230769 0.9245283    0
##
## Spec
##           Min.    1st Qu.    Median      Mean    3rd Qu.      Max. NA's
## ctree  0.6470588 0.6519608 0.7121212 0.7220143 0.7803030 0.8484848    0
## gbmA    0.6176471 0.6742424 0.7205882 0.7397504 0.8181818 0.8484848    0
```

Based on the cross-validation results on train data, bootstrap has a higher mean ROC value, implies bootstrap method performs better than classification tree.